

# Sistemas Distribuídos

José Orlando Pereira

Departamento de Informática  
Universidade do Minho



# Locks vs Variables

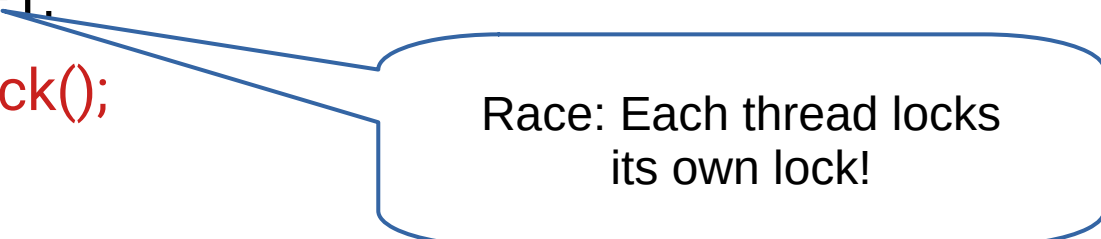
- “Which lock corresponds to each data item?”
- Multiple threads accessing some data item concurrently must have acquired the same lock
- Not automatic / not checked
- It is up to the developer to ensure this!

# Pitfall: Automatic variables

- Variables in methods are created on every invocation
- The method might still access shared state:
  - Instance variables
  - Class (static) variables

# Wrong

```
class SomeClass {  
    int s;  
    void doSomething() {  
        Lock l = new ReentrantLock();  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```



Race: Each thread locks  
its own lock!

# Solution

```
class SomeClass {  
    int s;  
    Lock l = new ReentrantLock();  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

Solution: Same scope for  
shared state and lock

# Pitfall: Class/global variables

- Variables marked with “static” in Java are global and (probably) need concurrency control
  - Not if marked “final”
  - Not if the class is used by a single thread

# Wrong

```
class SomeClass {  
    private static int s;  
    void doSomething() {  
        s = s+1;  
    }  
}
```



Race!

# Still wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private static int s;  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

There is one lock for each object, but s is shared!



# Solution

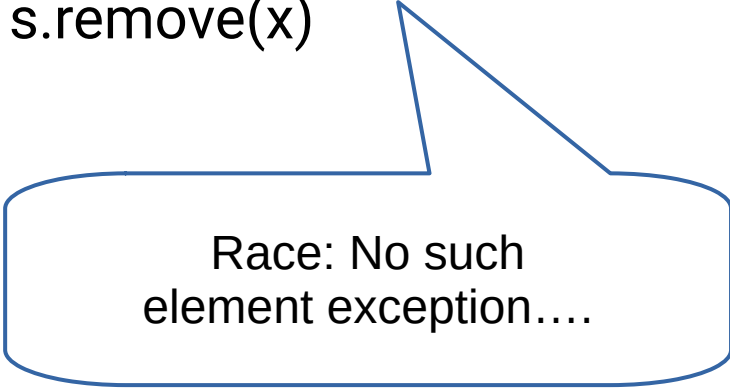
```
class SomeClass {  
    private static Lock l = new ReentrantLock();  
    private static int s;  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

# Pitfall: Encapsulated locks

- Keep variables and the corresponding lock encapsulated within the same object
- (The default using old-style “synchronized” in Java.)

# Wrong

```
class SomeClass {  
    SomeState s;  
    void doSomething() {  
        if (s.contains(x))  
            s.remove(x)  
    }  
}
```



Race: No such  
element exception....

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

# Solution

```
class SomeClass {  
    private Lock l;  
    SomeState s;  
    void doSomething() {  
        l.lock();  
        if (s.contains(x))  
            s.remove(x)  
        l.unlock();  
    }  
}
```

Now useless...

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

# Better solution

```
class SomeClass {  
    private Lock l;  
    SomeState s;  
    void doSomething() {  
        l.lock();  
        if (s.contains(x))  
            s.remove(x)  
        l.unlock();  
    }  
}
```

```
class SomeState {  
private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Rely on locking by the callers. This is done by Java Collections (Lists, ...)


# Pitfall: Shared vs thread-local state

- Program state often contains:
  - Local thread state in workers
  - Shared state, used by all threads
- Both are objects, with instance variables

# Wrong

```
class Worker
    extends Thread {
    Lock l = new ...;
    SharedState s;
    void doSomething() {
        l.lock(); s.doit(); l.unlock();
    }
    public void run() {
        doSomething();
    }
}
```


```
class SharedState {
    public void doit() {
        ...
    }
}
```



# Solution

```
class Worker
    extends Thread {
    SharedState s;
    void doit() {
        s.doit();
    }
    public void run() {
        doit();
    }
}
```

```
class SharedState {
    Lock l = new ...;
    public void doit() {
        l.lock(); ... l.unlock();
    }
}
```





# Quiz


- Two variables:
  - `int i=0, j=0;`
- Writer code:
  - `i=1; j=1;`
- Reader code:
  - `rj=j; ri=i; System.out.println(rj+", "+ri);`
- Possible results:
  - a) 0, 0 ✓
  - b) 1, 1 ✓
  - c) 0, 1 ✓
  - d) 1, 0 ✓

→ Why!?!?

} running  
concurrently!

# Memory order

- Steps to write a variable:
  1. Write to cache
  2. Flush cache to memory



(Image from <http://booksite.elsevier.com/9780123705914/?ISBN=9780123705914> . CC By-SA-3.0.)

# Memory order

- Possible outcome with two variables:


1. Write i to cache

2. Write j to cache

3. Flush j from cache to memory

5. Flush i from cache to memory

4. Paradox observed if i,j read here!!



(Image from <http://booksite.elsevier.com/9780123705914/?ISBN=9780123705914> . CC By-SA-3.0.)


# Pitfall: Read races

```
class X {  
    private Y y;  
    void changeY() {  
        l.lock();  
        tmp.i = 1;  
        y = tmp;  
        l.unlock();  
    }  
    int getY() {return y.i;}  
}
```

- Can we omit synchronization in getters?

# Pitfall: Read races

```
class X {  
    private Y y;  
    void changeY() {  
        l.lock();  
        tmp.i = 1;  
        y = tmp;  
        l.unlock();  
    }  
    int getI() { return y.i;}  
}
```



- Can we omit synchronization in getters?
  - **NO!**
- Can read inconsistent Y fields!
- In this case:
  - reader might not see `y.i == 1!!!!`

# Pitfall: Collections and getters

- Getter methods may return references to shared collections (or other mutable objects)
  - Iterators include references to the original object!

# Wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    List getElements() {  
        try { l.lock();  
            return l;  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
List l = l.getElements();  
l.add(...);
```



Race!

# Wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    Iterator getElements() {  
        try { l.lock();  
            return l.iterator();  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
    Iterator i = l.getElements();  
    while(i.hasNext())
```

...



Race!



# Still wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    List getElements() {  
        try { l.lock();  
            return l.clone();  
        } finally { l.unlock(); }  
    }  
}
```


```
SomeClass s = ...;  
List l = l.getElements();  
l.add(...);
```

Not adding to  
the list...

**...reconsider encapsulated lock!**

# Summary

- There is no simple rule to match locks with variables
- Some thinking needed... :-)



# Scaling up

- Example:
  - In a distributed database table with millions of records
  - Executing “select sum(x) from ... where ...” queries
  - Updating records
- Do we use a single lock?
  - Cannot run more than one query at the same time
- Do we use a lock for each line?
  - Way too many individual locks!

# Readers-Writers locks

- Strict mutual exclusion with locks is too conservative:
  - More than one reader would not be a problem
  - A writer must exclude all others (readers and writers)
- Different methods for readers and writers:

```
interface ReadWriteLock {  
    Lock readLock();  
    Lock writeLock();  
}
```
- More costly than a simple lock



# Readers-Writers locks

- R/W locks also known as shared locks in database management systems:
  - Readers lock  $\Leftrightarrow$  Shared mode
  - Writers lock  $\Leftrightarrow$  eXclusive mode
- Behavior described by a compatibility matrix:

Mode	R/S	W/X
R/S	Yes	No
W/X	No	No

# Readers-Writers example

```
int v;  
Lock l = new  
    ReentrantLock();  
void doSomething() {  
    l.lock();  
    v++;  
    l.unlock();  
}  
int getV() {  
    try { l.lock();  
        return v;  
    } finally { l.unlock(); }  
}
```

```
int v;  
ReadWriteLock l = new  
    ReentrantReadWriteLock();  
void doSomething() {  
    l.writeLock().lock();  
    v++;  
    l.writeLock().unlock();  
}  
int getV() {  
    try { l.readLock().lock();  
        return v;  
    } finally { l.readLock().unlock(); }  
}
```

... not worth it for such simple operations!

# Revisiting collections with 2PL+RW lock

- ```
void shoot(String sn, String tn) {  
    l.readLock().lock();  
    Player s = players.get(sn);  
    Player t = players.get(tn);  
    Stream.of(sn,tn).sorted()  
        .forEach(n→players.get(n).l.lock());  
    l.readLock().unlock();  
    t.life--;  
    t.l.unlock();  
    s.score++;  
    s.l.unlock();  
}
```

Allow multiple  
threads to acquire  
locks concurrently

Sorting is needed  
again

# Readers-Writers fairness

- Priority to readers
  - Allow more readers in, even if a writer is waiting
  - The writer may starve...
- Priority to writers
  - Do not allow more readers in, if a writer is waiting
  - Less concurrency among readers